

External-Memory Sorting

- External-memory algorithms
 - When data do not fit in main-memory
- External-memory sorting
 - Rough idea: sort peaces that fit in main-memory and “merge” them
- Main-memory merge sort:
 - The main part of the algorithm is Merge
 - Let's merge:
 - 3, 6, 7, 11, 13
 - 1, 5, 8, 9, 10

Main-Memory Merge Sort

Merge-Sort(A)

01 **if** length(A) > 1 **then**

02 Copy the first half of A into array A1

03 Copy the second half of A into array A2

04 **Merge-Sort**(A1)

05 **Merge-Sort**(A2)

06 **Merge**(A, A1, A2)


Divide

Conquer

Combine

■ Running time?

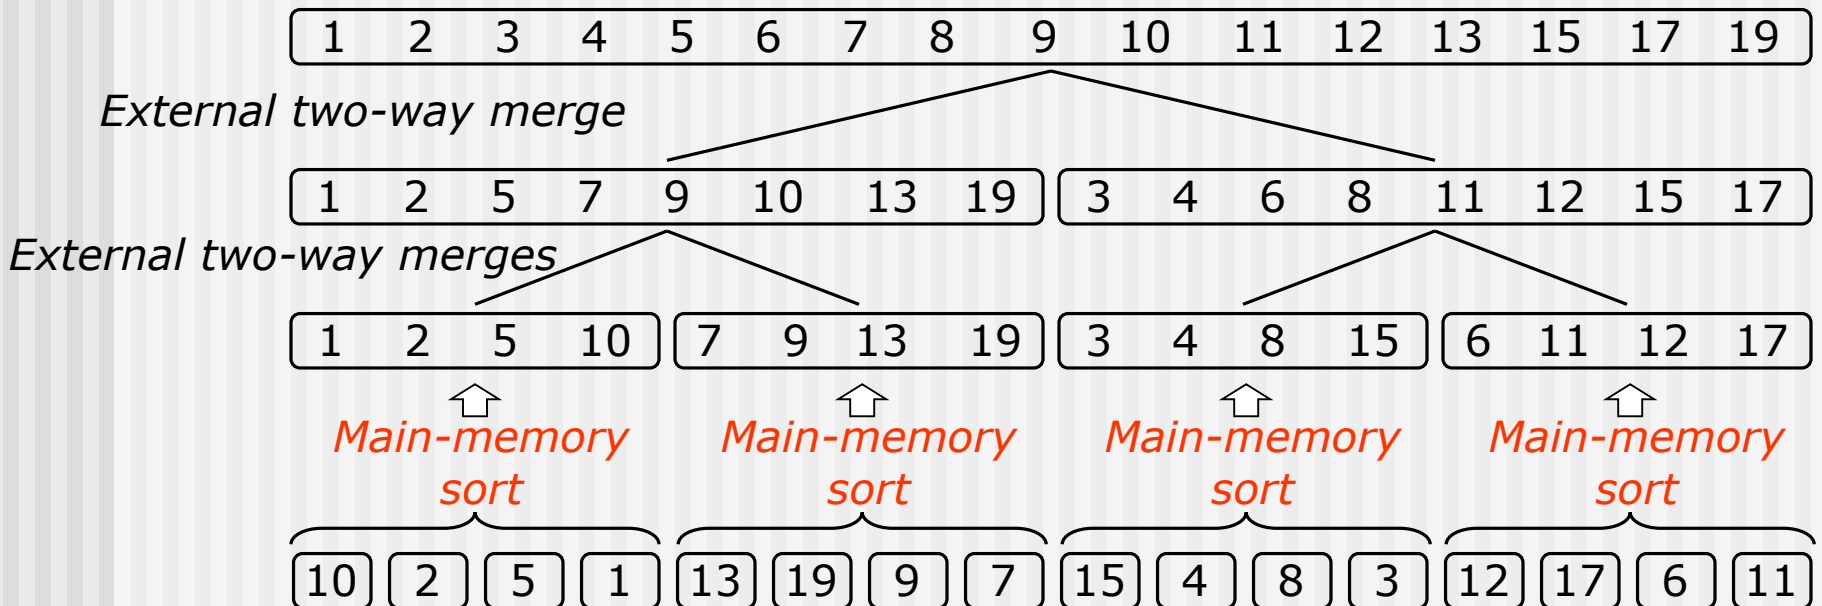
Merge-Sort Recursion Tree



- In each level: merge *runs* (sorted sequences) of size x into runs of size $2x$, decrease the number of runs twofold.
- What would it mean to run this on a file in external memory?

External-Memory Merge-Sort

- Idea: **increase the size of initial runs!**
 - Initial runs – the size of available main memory (M data elements)

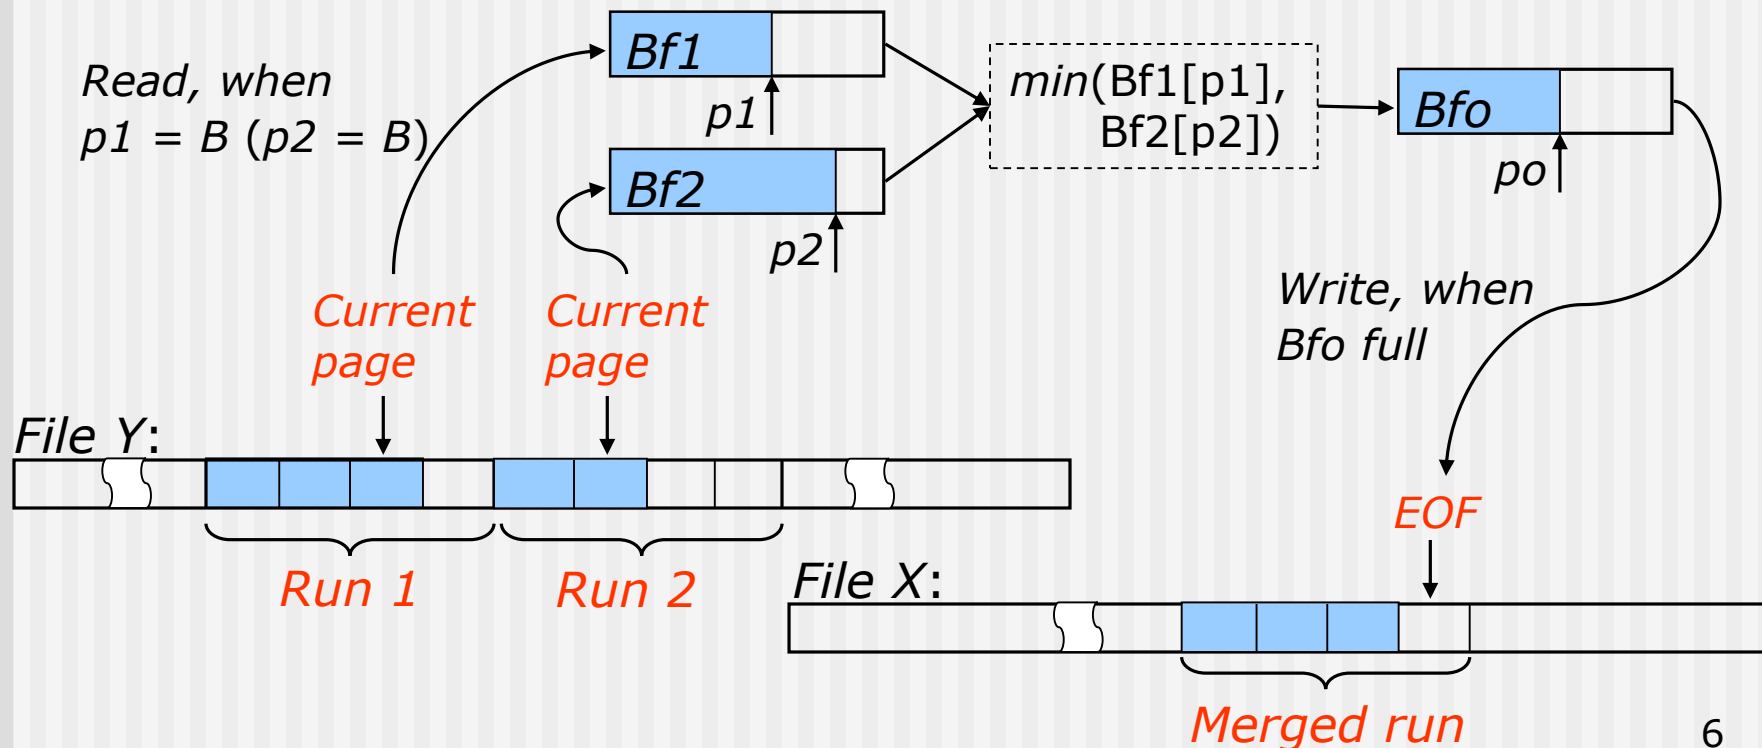


External-Memory Merge Sort

- Input file X , empty file Y
- *Phase 1: Repeat until end of file X :*
 - Read the next M elements from X
 - Sort them in main-memory
 - Write them at the end of file Y
- *Phase 2: Repeat while there is more than one run in Y :*
 - Empty X
 - $MergeAllRuns(Y, X)$
 - X is now called Y , Y is now called X

External-Memory Merging

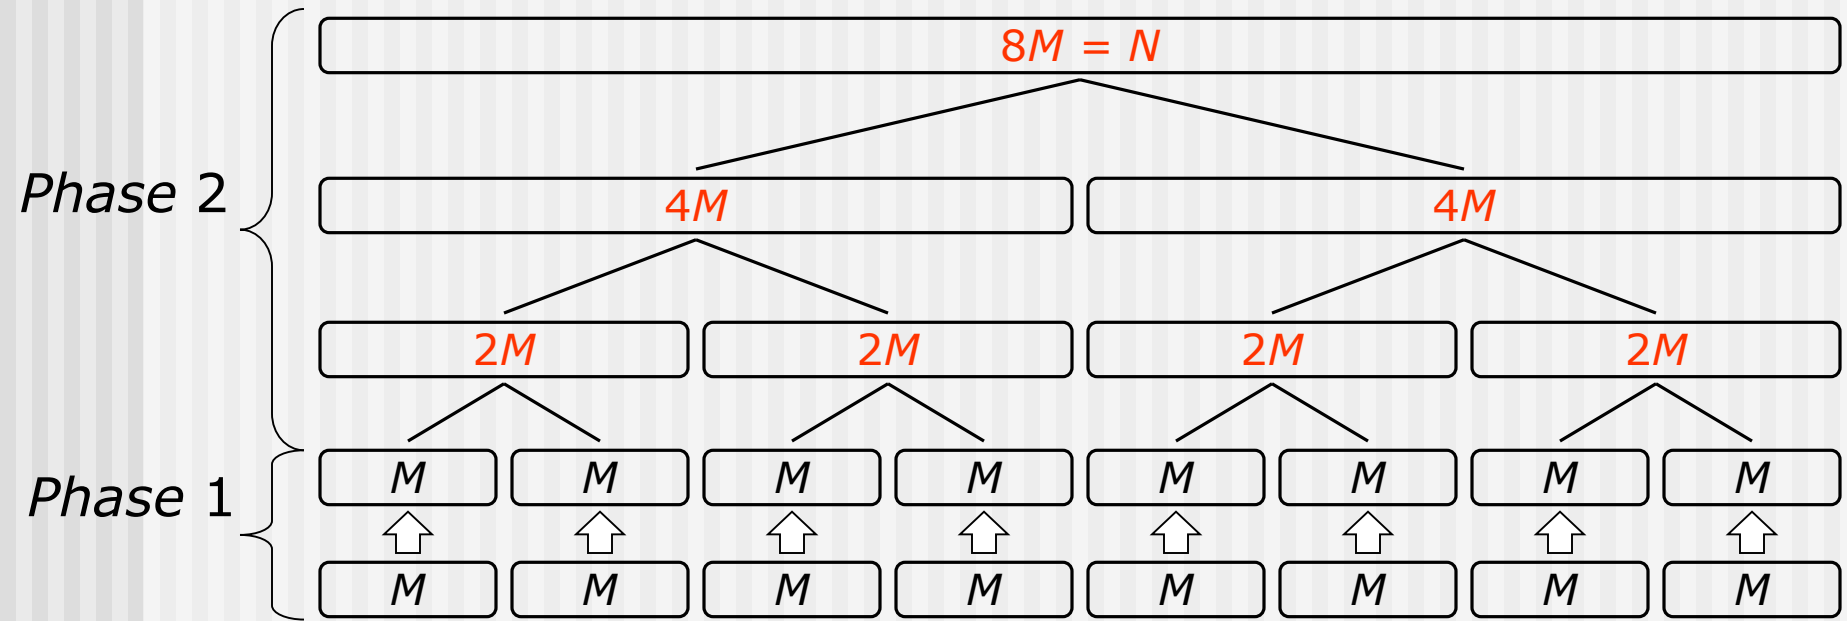
- *MergeAllRuns(Y, X)*: repeat until the end of *Y*:
 - Call *TwowayMerge* to merge the next two runs from *Y* into one run, which is written at the end of *X*
- *TwowayMerge*: uses three main-memory arrays of size *B*



Analysis: Assumptions

- Assumptions and notation:
 - Disk page size:
 - B data elements
 - Data file size:
 - N elements, $n = N/B$ disk pages
 - Available main memory:
 - M elements, $m = M/B$ pages

Analysis




- **Phase 1:**
 - Read file X, write file Y: $2n = O(n)$ I/Os
- **Phase 2:**
 - One iteration: Read file Y, write file X: $2n = O(n)$ I/Os
 - Number of iterations: $\log_2 N/M = \log_2 n/m$

Analysis: Conclusions

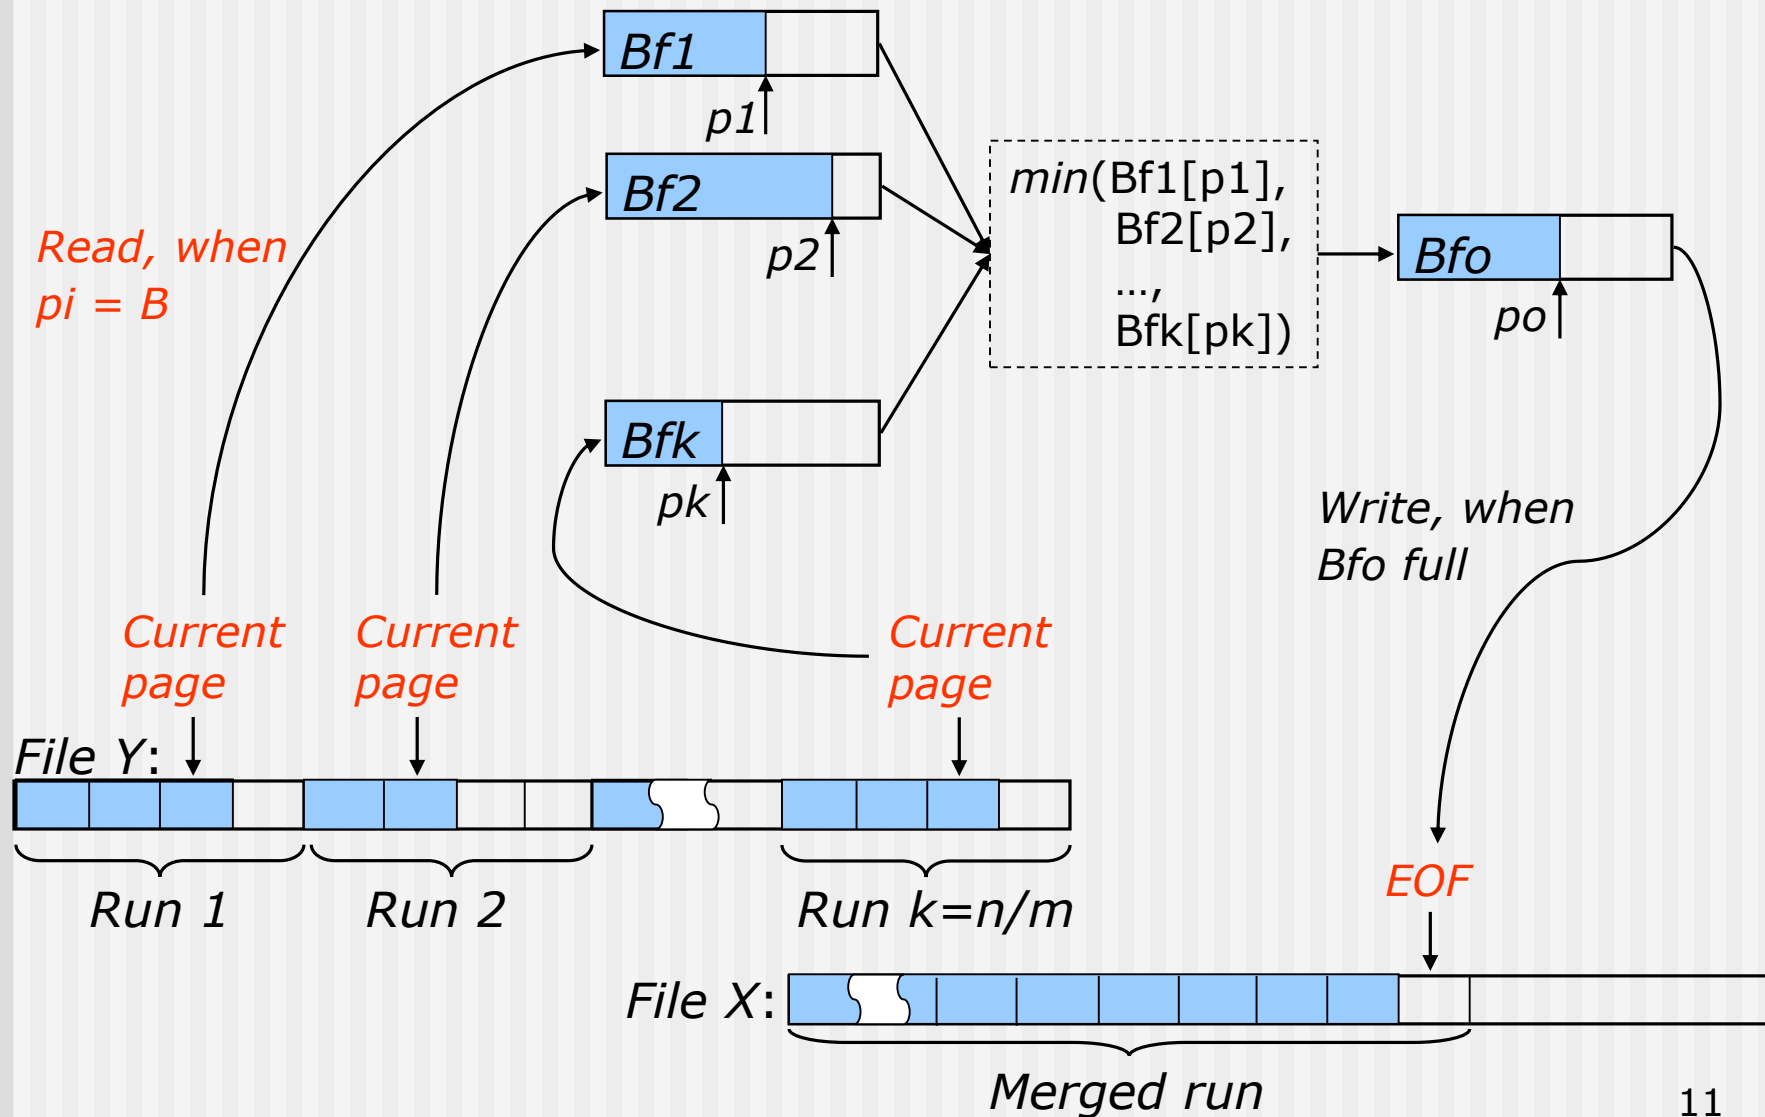
- Total running time of external-memory merge sort: $O(n \log_2 n/m)$
- We can do better!
- Observation:
 - Phase 1 uses all available memory
 - Phase 2 uses just 3 pages out of m available!!!

Two-Phase, Multiway Merge Sort

- Idea: merge all runs at once!
 - Phase 1: the same (do internal sorts)
 - Phase 2: perform *MultiwayMerge*(Y, X)



Multiway Merging



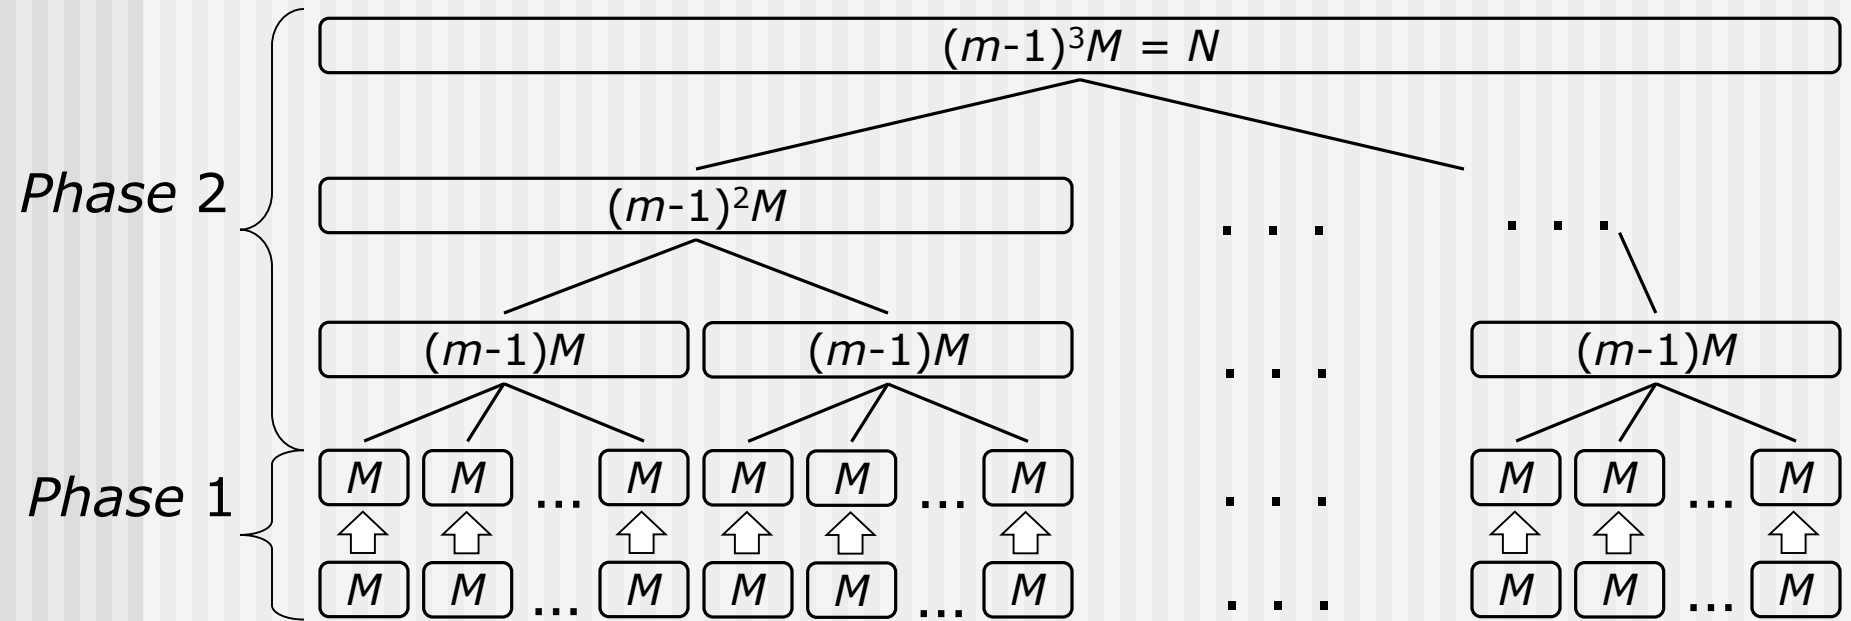
Analysis of TPMMS

- Phase 1: $O(n)$, Phase 2: $O(n)$
- Total: $O(n)$ I/Os!
- The catch: files only of “limited” size can be sorted
 - Phase 2 can merge a maximum of $m-1$ runs.
 - Which means: $N/M < m-1$
 - *How large files can we sort with TPMMS on a machine with 128Mb main memory and disk page size of 16Kb?*

General Multiway Merge Sort

- What if a file is very large or memory is small?
- General *multiway merge sort*:
 - **Phase 1**: the same (do internal sorts)
 - **Phase 2**: do as many iterations of merging as necessary until only one run remains
Each iteration repeatedly calls *MultiwayMerge*(Y, X) to merge groups of $m-1$ runs until the end of file Y is reached.

Analysis



- **Phase 1:** $O(n)$, each iteration of phase 2: $O(n)$
- How many iterations are there in phase 2?
 - Number of iterations: $\log_{m-1} N/M = \log_m n$
- Total running time: $O(n \log_m n)$ I/Os

Conclusions

- External sorting can be done in $O(n \log_m n)$ I/O operations for any n
 - This is asymptotically optimal
- In practice, we can usually sort in $O(n)$ I/Os
 - Use two-phase, multiway merge-sort